

CS 559: Machine Learning Fundamentals and Applications

Lecture 4

Lecturer: Xinchao Wang

xinchao.wang@stevens.edu

Teaching Assistant: Yiding Yang

yyang99@stevens.edu

Overview

- Parameter Estimation
 - Frequentist or Maximum Likelihood approach (cont.)
 - Bayesian approach (Barber Ch. 8 and DHS Ch. 3)
- Cross-validation
- Overfitting
- Non-parametric Techniques

MLE Classifier Example

Data

- Pima Indians Diabetes Database

- <http://archive.ics.uci.edu/ml/datasets/Pima+Indians+Diabetes>
- Number of Instances: 768
- Number of Attributes: 8 plus class
- Class Distribution: (class value 1 is interpreted as "tested positive for diabetes")

- | Class Value | Number of instances |
|-------------|---------------------|
| 0 | 500 |
| 1 | 268 |

Data

Attributes: (all numeric-valued)

1. Number of times pregnant
2. Plasma glucose concentration a 2 hours in an oral glucose tolerance test
3. Diastolic blood pressure (mm Hg)
4. Triceps skin fold thickness (mm)
5. 2-Hour serum insulin (mu U/ml)
6. Body mass index (weight in kg/(height in m)²)
7. Diabetes pedigree function
8. Age (years)
9. Class variable (0 or 1)

Simple MLE Classifier

```
data = dlmread('pima-indians-diabetes.data');
```

```
data = reshape(data, [], 9);
```

```
% use randperm to re-order data.
```

```
% ignore if not using Matlab
```

```
rp = randperm(length(data));
```

```
data=data(rp,:);
```

```
train_data = data(1:length(data)/2,:);
```

```
test_data = data(length(data)/2+1:end,:);
```

```
% pick a feature
active_feat = 3;

% training
mean1 =
    mean(train_data(train_data(:,9)==0,active_feat))
mean2 =
    mean(train_data(train_data(:,9)==1,active_feat))

var1 = var(train_data(train_data(:,9)==0,active_feat))
var2 = var(train_data(train_data(:,9)==1,active_feat))

prior1tmp = length(train_data(train_data(:,9)==0));
prior2tmp = length(train_data(train_data(:,9)==1));

prior1 = prior1tmp/(prior1tmp+prior2tmp)
prior2 = prior2tmp/(prior1tmp+prior2tmp)
```

```
% testing
correct=0;
wrong=0;
```

$$f(x \mid \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

```
for i=1:length(test_data)
    lklhood1 = exp(-(test_data(i,active_feat)-mean1)^2/(2*var1))
    /sqrt(var1);
    lklhood2 = exp(-(test_data(i,active_feat)-mean2)^2/(2*var2));
    /sqrt(var2);

    post1 = lklhood1*prior1;
    post2 = lklhood2*prior2;

    if(post1 > post2 && test_data(i,9) == 0)
        correct = correct+1;
    elseif(post1 < post2 && test_data(i,9) == 1)
        correct = correct+1;
    else
        wrong = wrong+1;
    end
end
end
```


Training/Test Split

- Randomly split dataset into two parts:
 - Training data
 - Test data
- Use training data to optimize parameters
- Evaluate error using test data

Training/Test Split

- How many points in each set?
- Very hard question
 - Too few points in training set, learned classifier is bad
 - Too few points in test set, classifier evaluation is insufficient
- Cross-validation
- Leave-one-out cross-validation
- Bootstrapping

Cross Validation





- In practice
- Available data => training and validation
- Train on the training data
- Test on the validation data
- k-fold cross validation:
 - Data randomly separated into k groups
 - Each time k-1 groups used for training and one as testing

Cross Validation and Test Accuracy

- If we select parameters so that CV is highest:
 - Does CV represent future test accuracy?
 - Slightly different
- So split available data with class labels, into:
 - training
 - validation
 - testing

Cross Validation and Test Accuracy

- Using CV on training + validation
- Classify test data with the best parameters from CV



Overfitting

- Prediction error: probability of test pattern not in class with max posterior (true)
- Training error: probability of test pattern not in class with max posterior (estimated)
- Classifier optimized w.r.t. training error
 - Training error: optimistically biased estimate of prediction error

Overfitting


Overfitting: a learning algorithm overfits the ***training data*** if it outputs a solution **\mathbf{w}** when another solution **\mathbf{w}'** exists such that:

$$\text{error}_{\text{train}}(\mathbf{w}) < \text{error}_{\text{train}}(\mathbf{w}')$$


AND

$$\text{error}_{\text{true}}(\mathbf{w}') < \text{error}_{\text{true}}(\mathbf{w})$$

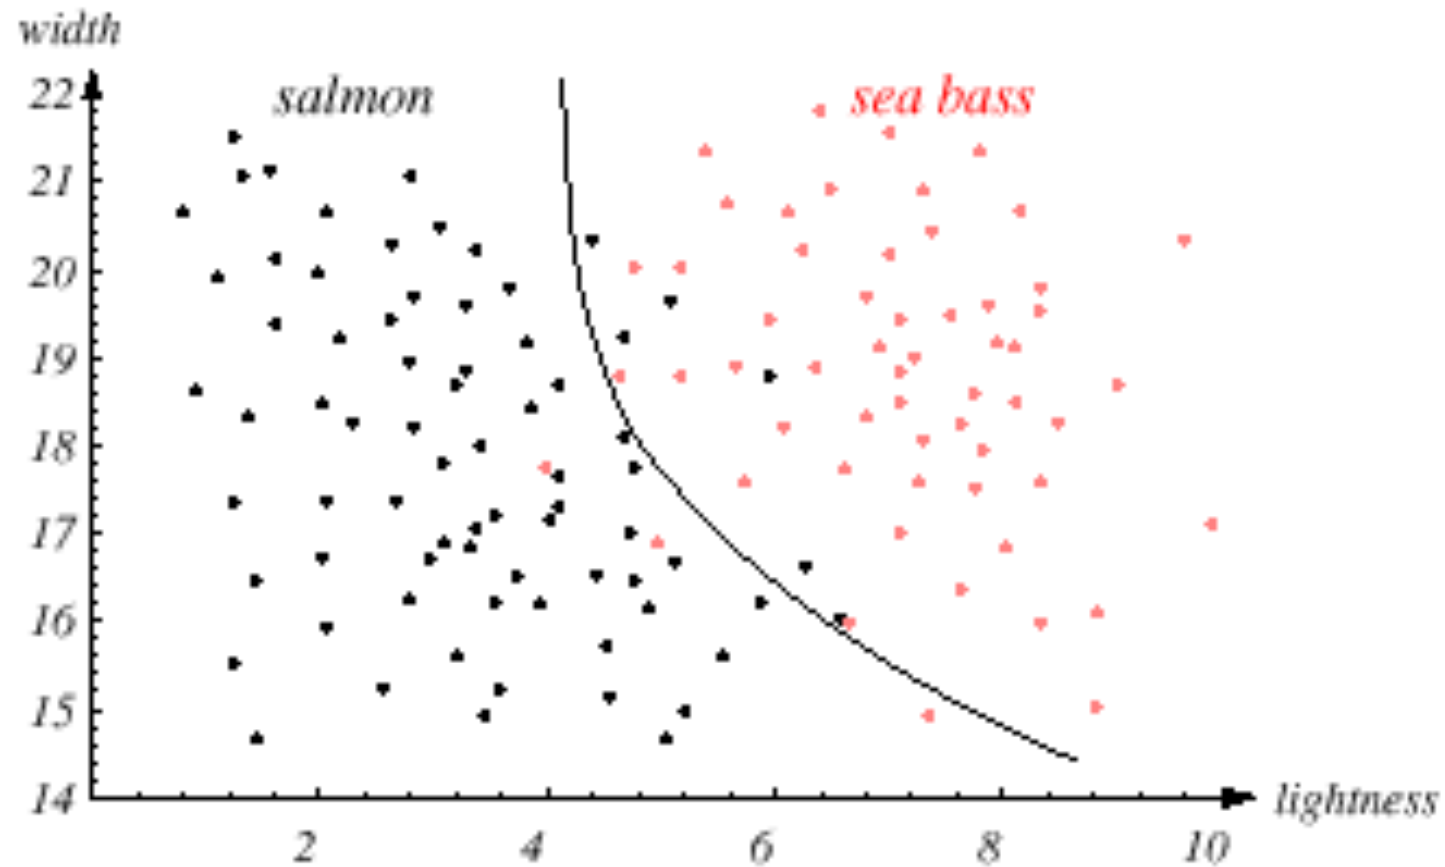
Fish Classifier from DHS Ch. 1




Minimum Training Error




Final Decision Boundary



Typical Behavior



Typical Behavior



Maximum a posteriori (MAP) Estimation

MLE

$$\begin{aligned}\theta_{MLE} &= \arg \max_{\theta} P(D|\theta) \\ &= \arg \max_{\theta} \log P(D|\theta) \\ &= \arg \max_{\theta} \log \Pi_i P(x_i|\theta) \\ &= \arg \max_{\theta} \sum_i \log P(x_i|\theta)\end{aligned}$$

MAP

- Recall Bayesian

$$P(\theta|D) = \frac{P(D|\theta)P(\theta)}{P(D)} \propto P(D|\theta)P(\theta)$$

- MAP


$$\begin{aligned}\theta_{MAP} &= \arg \max_{\theta} P(\theta|D) \\ &= \arg \max_{\theta} \log P(D|\theta) P(\theta) \\ &= \arg \max_{\theta} \log \Pi_i P(x_i|\theta) P(\theta) \\ &= \arg \max_{\theta} \Sigma_i \log P(x_i|\theta) P(\theta)\end{aligned}$$

MLE and MAP

$$\begin{aligned}\theta_{MLE} &= \arg \max_{\theta} P(D|\theta) \\ &= \arg \max_{\theta} \log P(D|\theta) \\ &= \arg \max_{\theta} \log \prod_i P(x_i|\theta) \\ &= \arg \max_{\theta} \sum_i \log P(x_i|\theta)\end{aligned}$$

$$\begin{aligned}\theta_{MAP} &= \arg \max_{\theta} P(\theta|D) \\ &= \arg \max_{\theta} \log P(D|\theta) P(\theta) \\ &= \arg \max_{\theta} \log \prod_i P(x_i|\theta) P(\theta) \\ &= \arg \max_{\theta} \sum_i \log P(x_i|\theta) P(\theta)\end{aligned}$$


MLE can be thought
as an uniform-prior
version of MAP!



Non-parametric Classification

The Histogram

- The simplest form of non-parametric density estimation is the histogram
 - Divide sample space in number of bins
 - Approximate the density at the center of each bin by the fraction of points that fall into the bin
 - Two parameters: bin width and starting position of first bin (or other equivalent pairs)
- Drawbacks:
 - Depends on position of bin centers
 - Often compute two histograms, offset by $\frac{1}{2}$ bin width
 - Discontinuities as an artifact of bin boundaries
 - Curse of dimensionality



Non-parametric Methods

- Non-parametric procedures can be used with **arbitrary distributions** and without the assumption that the forms of the underlying densities are known

Density Estimation

- Probability that a vector x will fall in region R is:

$$P = \int_{\mathfrak{R}} p(x') dx' \quad (1)$$

- P is a smoothed (or averaged) version of the density function $p(x)$ if we have a sample of size n ; therefore, the probability that k points fall in R is:

$$P_k = \binom{n}{k} P^k (1-P)^{n-k} \quad (2)$$

and the expected value for k is:

$$E(k) = nP \quad (3)$$

Answer the question: What is k , given n and P ?

ML Estimate

ML estimation of $P = \theta$

$\text{Max}_{\theta}(P_k | \theta)$ is reached for $\hat{\theta} = \frac{k}{n} \equiv P$

Therefore, the ratio k/n is a good estimate for the probability P and hence for the density function $p(\mathbf{x})$ (for large n)

Answer the question: What is P given k and n ?

The k-Nearest-Neighbor Rule

- **Goal:** Classify x by assigning it the label most frequently represented among the k nearest samples
- Use a voting scheme




FIGURE 4.15. The k -nearest-neighbor query starts at the test point \mathbf{x} and grows a spherical region until it encloses k training samples, and it labels the test point by a majority vote of these samples. In this $k = 5$ case, the test point \mathbf{x} would be labeled the category of the black points. From: Richard O. Duda, Peter E. Hart, and David G. Stork, *Pattern Classification*. Copyright © 2001 by John Wiley & Sons, Inc.

Matlab Example

```
data = dlmread('pima-indians-diabetes.data');
```

```
data = reshape(data,[],9);
```

```
% use randperm to re-order data. ignore if not using Matlab
```

```
rp = randperm(length(data));
```

```
data=data(rp,:);
```

```
%split = length(data)/2;
```

```
split = 300;
```

```
train_data = data(1:split,:);
```

```
test_data = data(split+1:end,:);
```



```
% pick features  
active_feat = [1:3];
```

```
% training  
% NOT NEEDED
```

```
% testing  
correct=0;  
wrong=0;
```

```

for i=1:length(test_data)

    sample=test_data(i,active_feat);

    dist = train_data(:,active_feat)-repmat(sample,length(train_data),1);
    dist = dist*dist';

    % we are only interested in the diagonal elements
    % DON'T USE QUADRATIC DISTANCE COMPUTATION IN PRACTICE
    fin_dist = diag(dist);
    [min_d index] = min(fin_dist);

    if(test_data(i,9) == train_data(index,9))
        correct = correct+1;
    else
        wrong = wrong+1;
    end
end
end

```

Supplement

Bayesian Parameter Estimation

- Gaussian Case
- General Estimation

Bayesian Estimation

- In MLE θ was assumed fixed
- In BE θ is a random variable
- Suppose we have some idea of the range where the parameters θ should be
 - Shouldn't we utilize this prior knowledge in hope that it will lead to better parameter estimation?

Bayesian Estimation

- Let θ be a random variable with prior distribution $P(\theta)$
 - This is the key difference between ML and Bayesian parameter estimation
 - This allows us to use a prior to express the uncertainty present before seeing the data
 - Frequentist approach does not account for uncertainty in θ

Motivation

- As in MLE, suppose $p(x|\theta)$ is completely specified if θ is given
- But now θ is a random variable with prior $p(\theta)$
 - Unlike MLE case, $p(x|\theta)$ is a conditional density
- After we observe the data D , using Bayes rule we can compute the posterior $p(\theta|D)$

Motivation

- Recall that for the MAP classifier we find the class ω_i that maximizes the posterior $p(\omega | D)$
- By analogy, a reasonable estimate of θ is the one that maximizes the posterior $p(\theta | D)$
- But θ is not our final goal, our final goal is the unknown $p(x)$
- Therefore a better thing to do is to maximize $p(x | D)$, this is as close as we can come to the unknown $p(x)$!


Parameter Distribution

- Assumptions:
 - $p(x)$ is unknown, but has known parametric form
 - Parameter vector θ is unknown
 - $p(x | \theta)$ is completely known
 - Prior density $p(\theta)$ is known
- Observation of samples provides posterior density $p(\theta | D)$
 - Hopefully peaked around true value of θ
- Treat each class separately and drop subscripts

- Converted problem of learning ***probability density*** function to learning ***parameter vector***
- Goal: compute $p(x | D)$ as best possible estimate of $p(x)$

$$p(x | D) = \int p(x, \theta | D) d\theta$$

$$p(x | D) = \int p(x | \theta, D) p(\theta | D) d\theta = \int p(x | \theta) p(\theta | D) d\theta$$


 $p(x)$ is completely known given θ ,
independent of samples in D

$$p(\mathbf{x} | D) = \int p(\mathbf{x} | \theta, D) p(\theta | D) d\theta = \int p(\mathbf{x} | \theta) p(\theta | D) d\theta$$

- Links class-conditional density $p(\mathbf{x} | D)$ to posterior density $p(\theta | D)$

Bayesian Parameter Estimation: Gaussian Case

Goal: Estimate θ using the a-posteriori density $P(\theta \mid D)$

- The univariate case: $p(\mu \mid D)$
 μ is the only unknown parameter

$$p(x \mid \mu) \sim N(\mu, \sigma^2)$$

$$p(\mu) \sim N(\mu_0, \sigma_0^2)$$

$$\begin{aligned} p(\mu | \mathbf{D}) &= \frac{p(\mathbf{D} | \mu) p(\mu)}{\int p(\mathbf{D} | \mu) p(\mu) d\mu} \\ &= \alpha \prod_{k=1}^{k=n} p(x_k | \mu) p(\mu) \end{aligned}$$

- α depends on \mathbf{D} , not μ
- Shows how training samples affect our idea about the true value of μ

$$p(\mu | \mathbf{D}) = \frac{p(\mathbf{D} | \mu) p(\mu)}{\int p(\mathbf{D} | \mu) p(\mu) d\mu} \quad (1)$$

$$= \alpha \prod_{k=1}^{k=n} p(x_k | \mu) p(\mu)$$

Reproducing density (remains Gaussian)

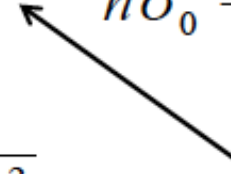
$$p(\mu | \mathbf{D}) \sim N(\mu_n, \sigma_n^2) \quad (2)$$

(1) and (2) yield:

$$\mu_n = \left(\frac{n\sigma_0^2}{n\sigma_0^2 + \sigma^2} \right) \hat{\mu}_n + \frac{\sigma^2}{n\sigma_0^2 + \sigma^2} \mu_0$$

and $\sigma_n^2 = \frac{\sigma_0^2 \sigma^2}{n\sigma_0^2 + \sigma^2}$


Empirical (sample) mean



$$\mu_n = \left(\frac{n\sigma_0^2}{n\sigma_0^2 + \sigma^2} \right) \hat{\mu}_n + \frac{\sigma^2}{n\sigma_0^2 + \sigma^2} \mu_0$$

$$\text{and } \sigma_n^2 = \frac{\sigma_0^2 \sigma^2}{n\sigma_0^2 + \sigma^2}$$

- μ is linear combination of empirical and prior information
- Each additional observation decreases uncertainty about μ



- The univariate case $p(x | D)$
 - $p(\mu | D)$ computed
 - $p(x | D)$ remains to be computed*

$$p(x | D) = \int p(x | \mu) p(\mu | D) d\mu \text{ is Gaussian}$$

It provides: $p(x | D) \sim N(\mu_n, \sigma^2 + \sigma_n^2)$

$$p(x | \mathbf{D}) \sim N(\mu_n, \sigma^2 + \sigma_n^2)$$

- We have:
 - Replaced mean with conditional mean
 - Increased variance to account for additional uncertainty in x due to inexact knowledge of mean

Bayesian Parameter Estimation: General Theory

- $p(x \mid D)$ computation can be applied to any situation in which the unknown density can be parameterized. The basic assumptions are:
 - The form of $p(x \mid \theta)$ is assumed known, but the value of θ is not known exactly
 - Our knowledge about θ is assumed to be contained in a known prior density $p(\theta)$
 - The rest of our knowledge θ is contained in a set D of n random variables x_1, x_2, \dots, x_n that follows $p(x)$

Recursive Bayes Learning

- Assume that training samples become available one by one

$$p(\mathbf{D}^n | \theta) = p(x_n | \theta) p(\mathbf{D}^{n-1} | \theta)$$

- Due to independence, result is independent of order:

$$p(\mathbf{D} | \theta) = \prod_{k=1}^{k=n} p(x_k | \theta)$$

The basic problem is:

“Compute the posterior density $p(\theta \mid D)$ ”
then “Derive $p(x \mid D)$ ”

Using Bayes formula, we have:

$$p(\theta \mid D) = \frac{p(D \mid \theta)p(\theta)}{\int p(D \mid \theta)p(\theta)d\theta}$$

And by the independence assumption:

$$p(D \mid \theta) = \prod_{k=1}^{k=n} p(x_k \mid \theta)$$

Estimation of $p(x | D)$

- The basic problem is: **Compute $p(x | D)$**

$$p(x | D) = \int \overset{\text{known}}{p(x | \theta)} \overset{\text{unknown}}{p(\theta | D)} d\theta$$

- Compute the posterior density $p(\theta | D)$

$$p(\theta | D) = \frac{p(D | \theta) p(\theta)}{\int p(D | \theta) p(\theta) d\theta}$$

- Then derive $p(x | D)$

ML vs. Bayesian Parameter Estimation: Summary

BE vs. MLE

- BE: $p(x|D)$ can be thought of as the weighted average of the proposed model for all possible values of θ

$$p(x|D) = \int \underbrace{p(x|\theta)}_{\text{proposed model with certain } \theta} \underbrace{p(\theta|D)}_{\text{support } \theta \text{ receives from the data}} d\theta$$

- Contrast this with the MLE solution which always gives us a single model:

$$p(x|\hat{\theta})$$

- When we have many possible solutions, taking their sum averaged by their probabilities seems better than pick just one solution

BE vs. MLE

- In practice, it may be hard to do integration analytically and we may have to resort to numerical methods
- The MLE solution requires differentiation, instead of integration, to get

$$p(x|\hat{\theta})$$

- Differentiation is easy and can always be done analytically

When do Maximum-Likelihood and Bayes Methods Differ?

- Equivalent asymptotically (for infinite training data)
 - For reasonable prior distributions
 - When prior $p(\theta)$ is uninformative and $p(\theta|D)$ is peaked
- MLE computationally cheaper, simpler solutions
- BE uses more information (more general model)

Naïve Bayes Classifier (not BE)

- Simple classifier that applies Bayes' rule with strong (naive) independence assumptions
- A.k.a. the "independent feature model"
- Often performs reasonably well despite simplicity