

# CS 559: Machine Learning Fundamentals and Applications


## Lecture 11

Lecturer: Xinchao Wang

[xinchao.wang@stevens.edu](mailto:xinchao.wang@stevens.edu)


Teaching Assistant: Yiding Yang

[yyang99@stevens.edu](mailto:yyang99@stevens.edu)




# Perceptrons

- Perceptrons
  - 1950s ~ 1960s, Frank Rosenblatt, inspired by earlier work by Warren McCulloch and Walter Pitts
- Standard model of artificial neurons



# Binary Perceptrons


- Inputs
  - Multiple binary inputs
- Parameters
  - Thresholds & weights
- Outputs
  - Thresholded weighted linear combination



$$\text{output} = \begin{cases} 0 & \text{if } \sum_j w_j x_j \leq \text{threshold} \\ 1 & \text{if } \sum_j w_j x_j > \text{threshold} \end{cases}$$

# Layered Perceptrons



- Layered, complex model
  - 1<sup>st</sup> layer, 2<sup>nd</sup> layer of perceptrons
- Perceptron rule
  - Weights, thresholds



$$\text{output} = \begin{cases} 0 & \text{if } w \cdot x + b \leq 0 \\ 1 & \text{if } w \cdot x + b > 0 \end{cases}$$

# Output Functions

- Sigmoid neurons
- Output  $\sigma(w \cdot x + b)$ ,  $\sigma(z) \equiv \frac{1}{1 + e^{-z}}$   
$$\frac{1}{1 + \exp(-\sum_j w_j x_j - b)}$$
- Sigmoid vs conventional thresholds



# Neural Nets for Computer Vision

Based on Tutorials at CVPR 2012 and 2014 by  
Marc'Aurelio Ranzato

# Key Ideas of Neural Nets

## **IDEA # 1**

Learn features from data

## **IDEA # 2**

Use differentiable functions that produce features efficiently

## **IDEA # 3**

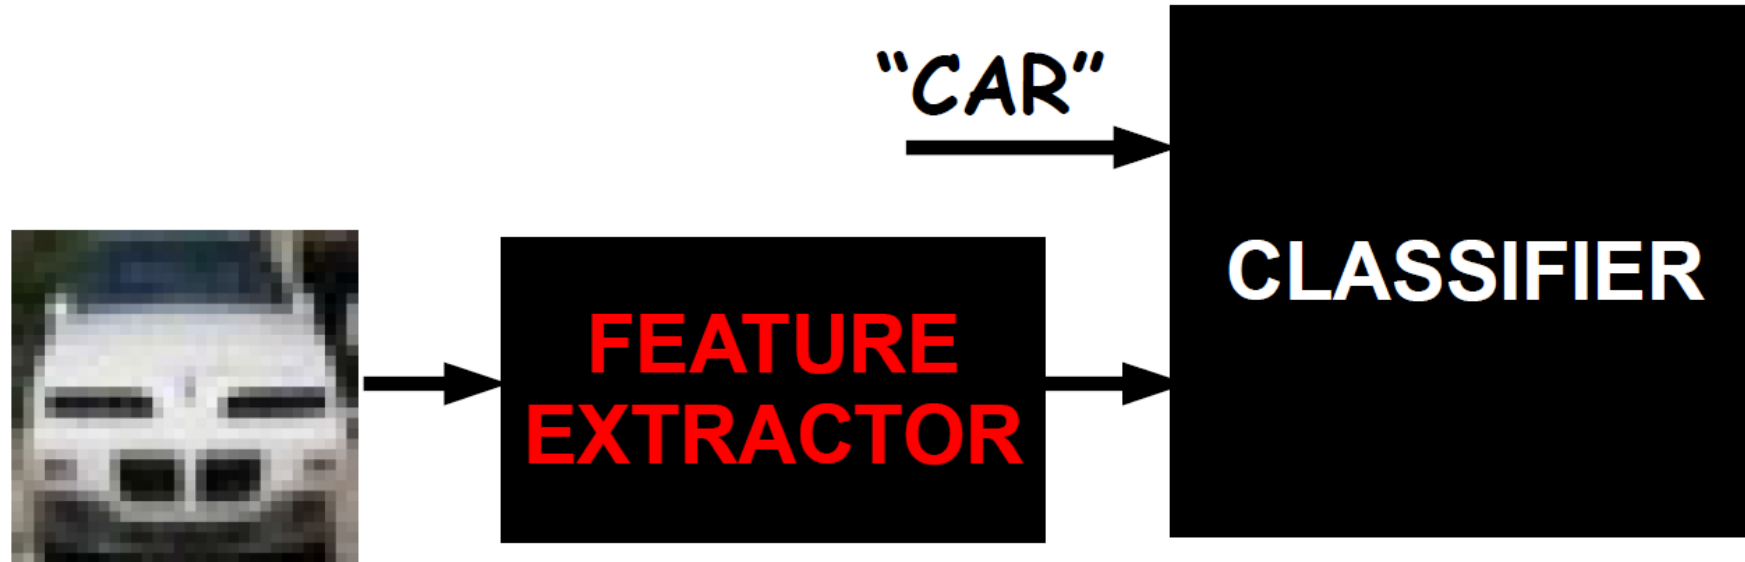
End-to-end learning:  
no distinction between feature extractor and classifier

## **IDEA # 4**

“Deep” architectures:  
cascade of simpler non-linear modules

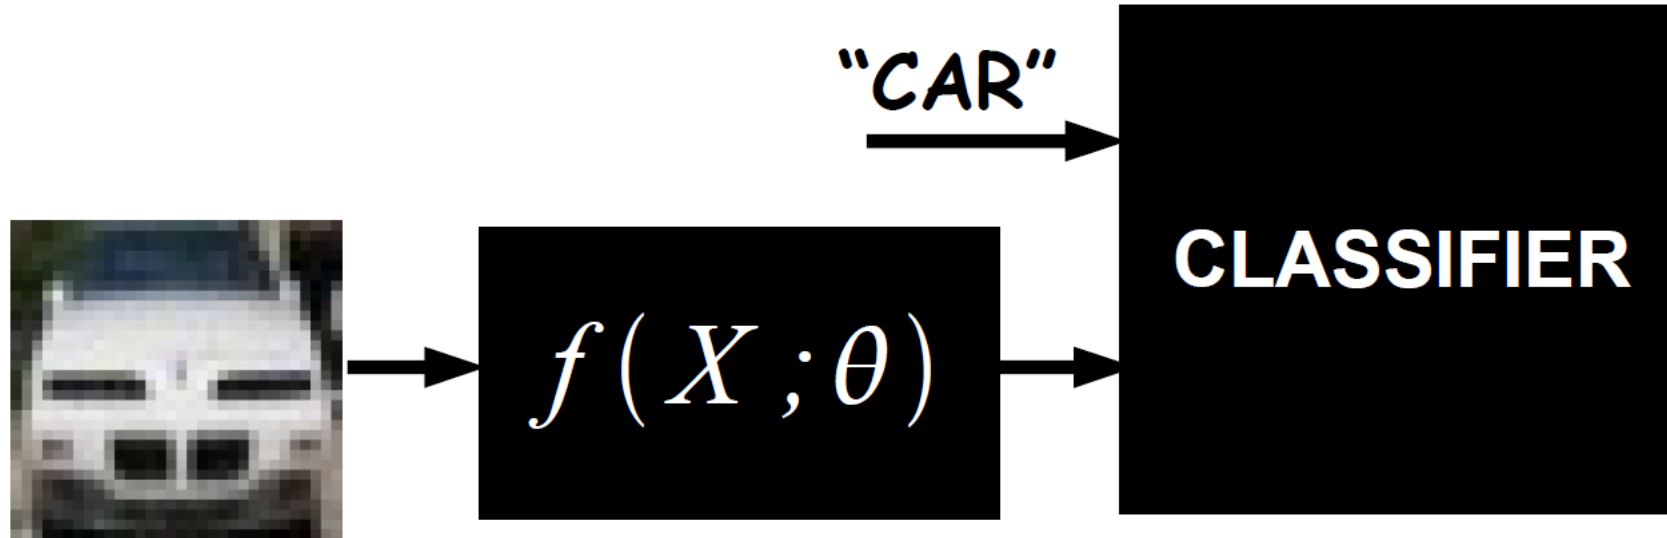


# Building an Object Recognition System



IDEA: Use data to optimize features for the given task


# Building an Object Recognition System



What we want: Use parameterized function such that


- a) features are computed efficiently
- b) features can be trained efficiently

# Building an Object Recognition System



- Everything becomes adaptive
- No distinction between feature extractor and classifier
- Big non-linear system trained from raw pixels to labels

# Building an Object Recognition System





**Q:** How can we build such a highly non-linear system?

**A:** By combining simple building blocks we can make more and more complex systems

# Building a Complicated Function

Simple Functions



One Example of  
Complicated Function





A red oval containing the complicated function:  $\log(\cos(\exp(\sin^3(x))))$

- Function composition is at the core of deep learning methods
- Each “simple function” will have parameters subject to training

# Implementing a Complicated Function

Complicated Function


$$\log(\cos(\exp(\sin^3(x))))$$



# Intuition Behind Deep Neural Nets




# Intuition Behind Deep Neural Nets



Each black box can have trainable parameters. Their composition makes a highly non-linear system.




# Intuition Behind Deep Neural Nets




System produces hierarchy of features


# Intuition Behind Deep Neural Nets



# Intuition Behind Deep Neural Nets



# Intuition Behind Deep Neural Nets



# Key Questions

- What is the input-output mapping?
- How are parameters trained?
- How computational expensive is it?
- How well does it work?

# Supervised Deep Learning

Marc'Aurelio Ranzato


# Supervised Learning

$\{(x_i, y_i), i=1 \dots P\}$  training set

$x_i$  i-th input training example

$y_i$  i-th target label

$P$  number of training examples



- Goal: predict the target label of unseen inputs

# Supervised Learning Examples


## Classification



→ “dog”



*classification*

## Denoising



*regression*

## OCR




→ “2 3 4 5”

*structured prediction*




# Supervised Deep Learning


## Classification



## Denoising



## OCR



# Neural Networks

Assumptions (for the next few slides):

- The input image is vectorized (disregard the spatial layout of pixels)
- The target label is discrete (classification)

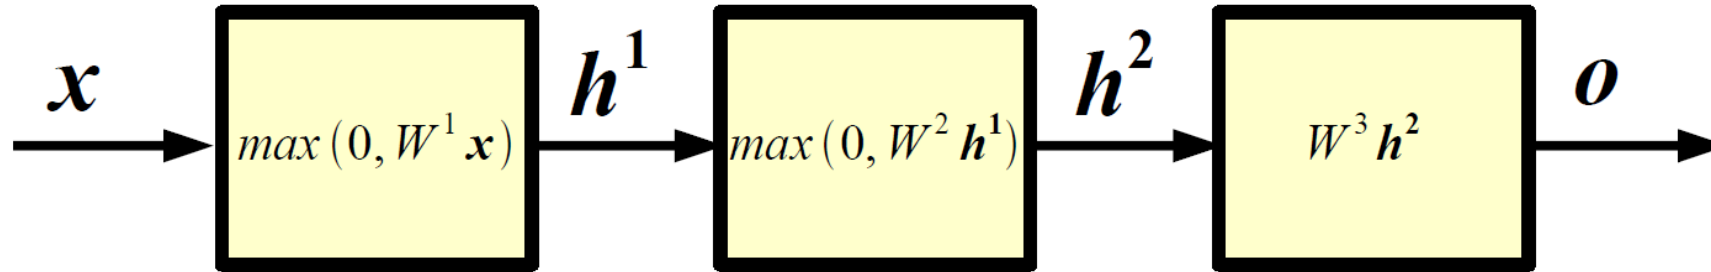
**Question:** what class of functions shall we consider to map the input into the output?

**Answer:** composition of simpler functions.

**Follow-up questions:** Why not a linear combination? What are the “simpler” functions? What is the interpretation?

**Answer:** later...

# Neural Networks: example



$x$  input

$h^1$  1-st layer hidden units

$h^2$  2-nd layer hidden units


$o$  output

Example of a 2 hidden layer neural network (or 4 layer network, counting also input and output)

# Forward Propagation

Forward propagation is the process of computing the output of the network given its input

# Forward Propagation




$$h^1 = \max(0, W^1 x + b^1)$$

$W^1$  1<sup>st</sup> layer weight matrix or weights

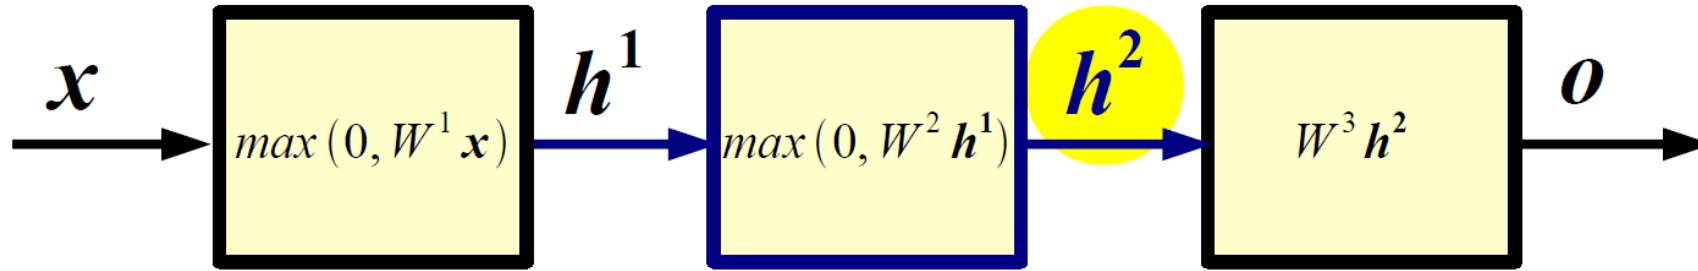
$b^1$  1<sup>st</sup> layer biases

- The non-linearity  $u = \max(0, v)$  is called **ReLU** in the DL literature.
- Each output hidden unit takes as input all the units at the previous layer: each such layer is called “**fully connected**”

# Rectified Linear Unit (ReLU)



# Forward Propagation



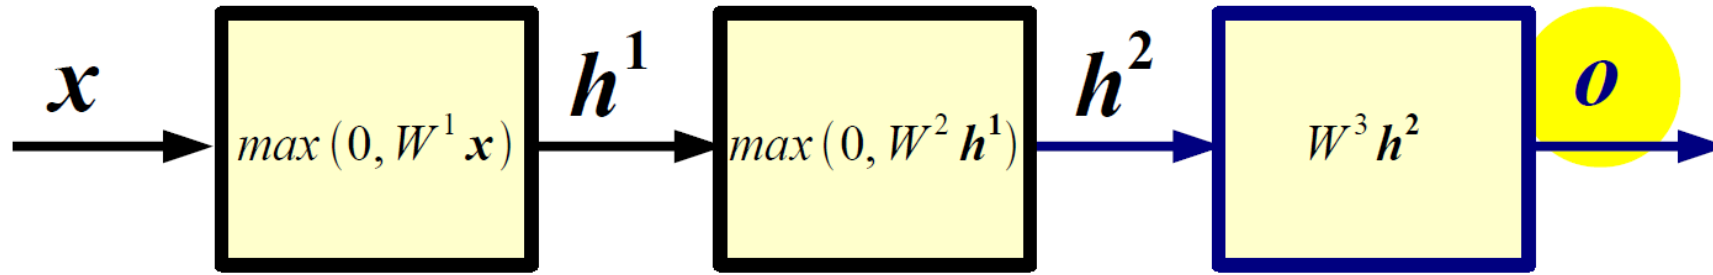
$$h^1 \in R^{N_1} \quad W^2 \in R^{N_2 \times N_1} \quad b^2 \in R^{N_2} \quad h^2 \in R^{N_2}$$

$$h^2 = \max(0, W^2 h^1 + b^2)$$

$W^2$  2<sup>nd</sup> layer weight matrix or weights

$b^2$  2<sup>nd</sup> layer biases

# Forward Propagation



$$h^2 \in R^{N_2} \quad W^3 \in R^{N_3 \times N_2} \quad b^3 \in R^{N_3} \quad o \in R^{N_3}$$

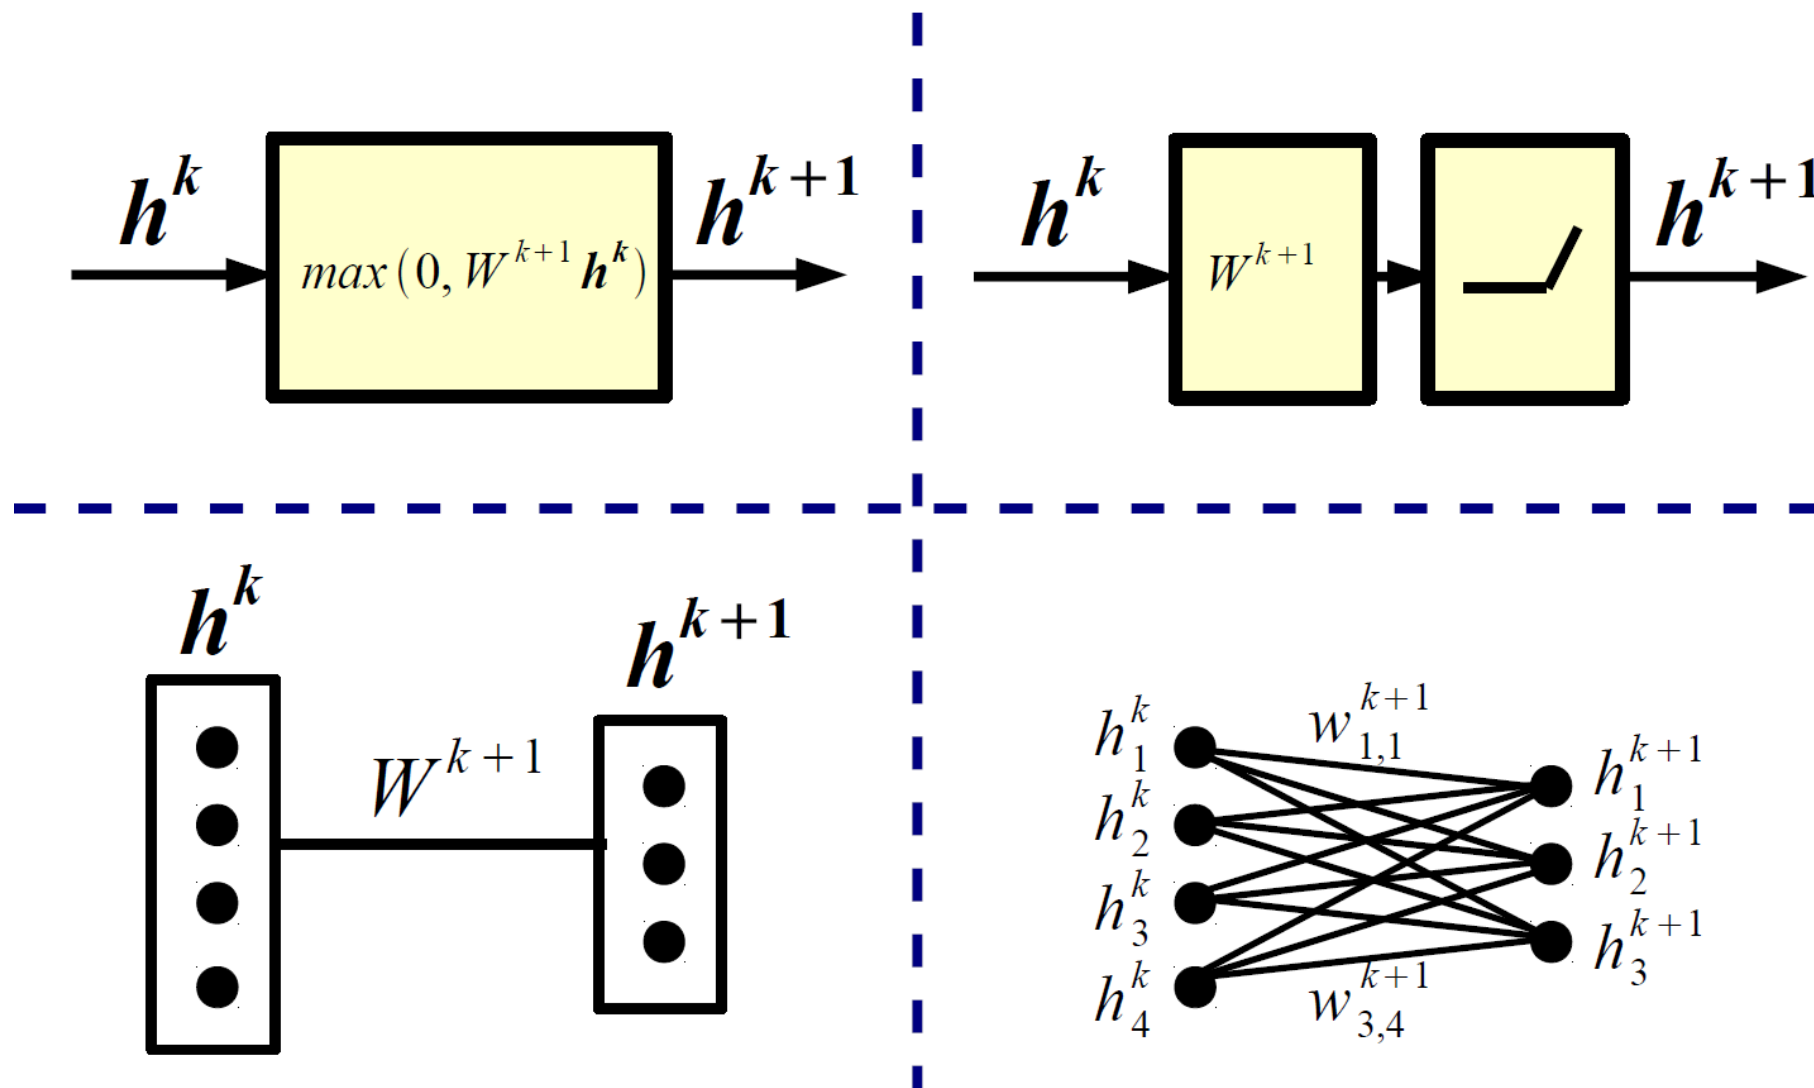
$$o = \max(0, W^3 h^2 + b^3)$$

$W^3$  3<sup>rd</sup> layer weight matrix or weights

$b^3$  3<sup>rd</sup> layer biases




# Alternative Graphical Representations



# Interpretation


- **Question:** Why can't the mapping between layers be linear?
- **Answer:** Because composition of linear functions is a linear function. Neural network would reduce to (1 layer) logistic regression.
- **Question:** What do ReLU layers accomplish?
- **Answer:** Piece-wise linear tiling: mapping is locally linear.



# Interpretation

- **Question:** Why do we need many layers?
- **Answer:** When input has hierarchical structure, the use of a hierarchical architecture is potentially more efficient because **intermediate computations** can be re-used.
- DL architectures are efficient also because they use **distributed representations** which are **shared** across classes.


[0 0 **1** 0 0 0 0 **1** 0 0 **1** **1** 0 0 **1** 0 ... ] truck feature



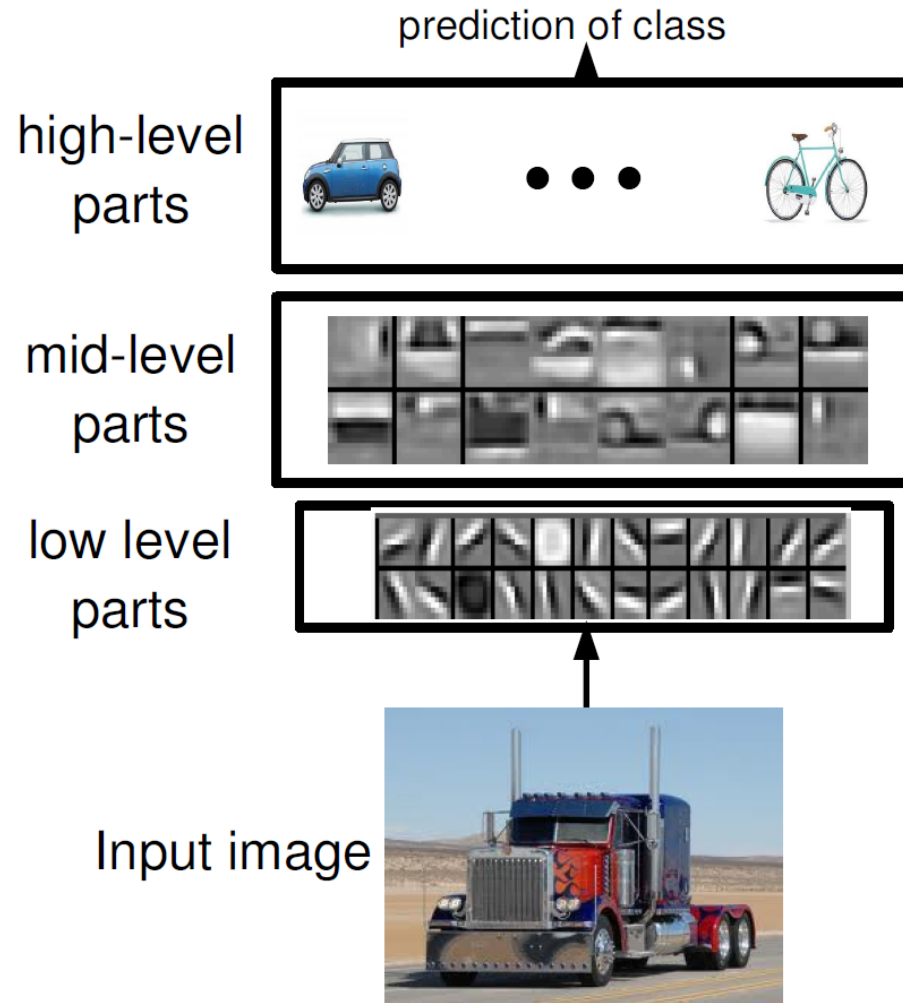
# Interpretation

[1 1 0 0 0 1 0 **1** 0 0 0 0 0 1 1 0 1...] motorbike

[0 0 1 0 0 0 0 **1** 0 0 1 1 0 0 1 0...] truck



# Interpretation



- Distributed representations
- Feature sharing
- Compositionality

# Interpretation

**Question:** What does a hidden unit do?

**Answer:** It can be thought of as a classifier or feature detector.

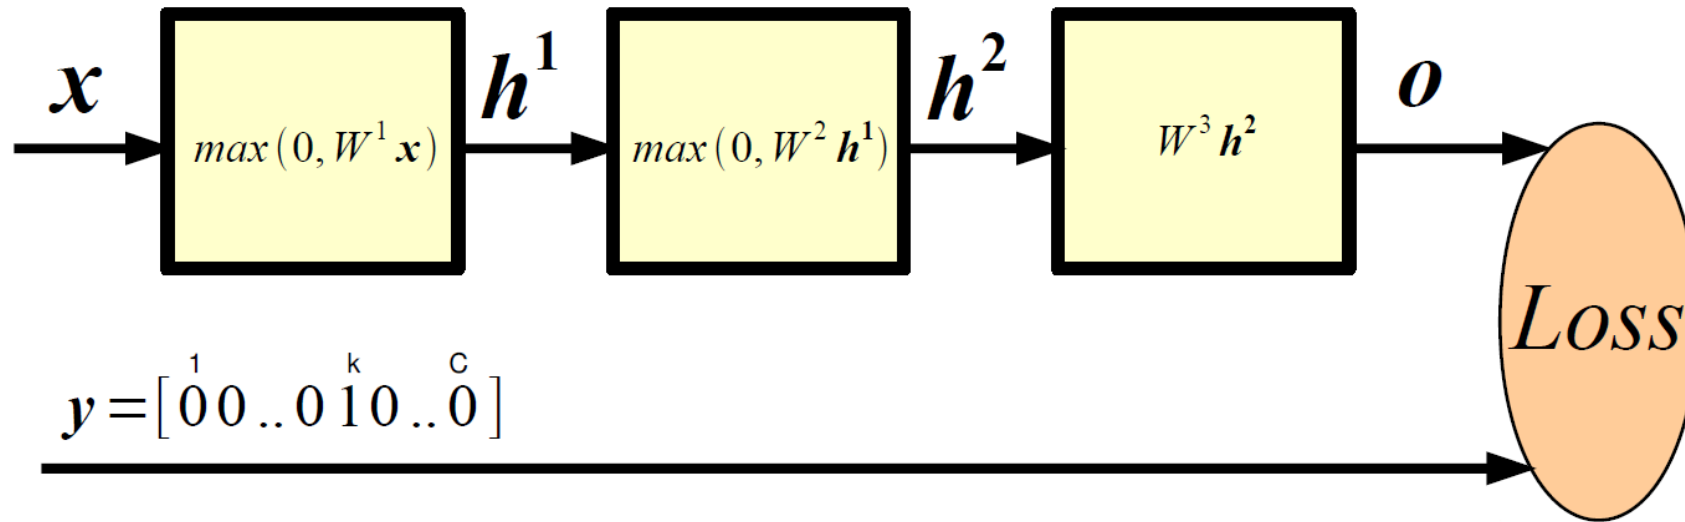
**Question:** How many layers? How many hidden units?

**Answer:** Cross-validation or hyper-parameter search methods are the answer. In general, the wider and the deeper the network the more complicated the mapping.

**Question:** How do I set the weight matrices?

**Answer:** Weight matrices and biases are learned. First, we need to define a measure of quality of the current mapping. Then, we need to define a procedure to adjust the parameters.

# How Good is a Network



- Probability of class  $k$  given input (softmax):

$$p(c_k = 1 | \mathbf{x}) = \frac{e^{o_k}}{\sum_{j=1}^C e^{o_j}}$$

- (Per-sample) **Loss**; e.g., negative log-likelihood (good for classification of small number of classes):

$$L(\mathbf{x}, \mathbf{y}; \boldsymbol{\theta}) = - \sum_j y_j \log p(c_j | \mathbf{x})$$

# Training

- Learning consists of minimizing the loss (plus some regularization term) w.r.t. parameters over the whole training set.

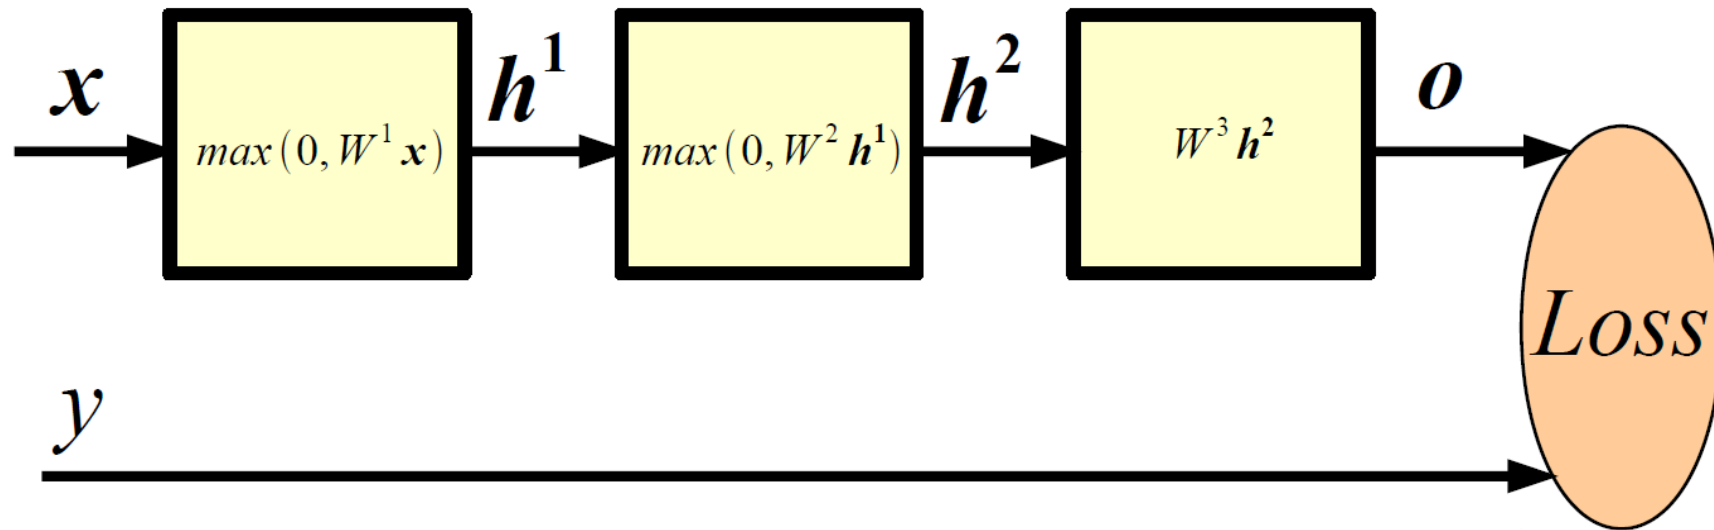
$$\boldsymbol{\theta}^* = \arg \min_{\boldsymbol{\theta}} \sum_{n=1}^P L(\boldsymbol{x}^n, y^n; \boldsymbol{\theta})$$

**Question:** How to minimize a complicated function of the parameters?

**Answer:** Chain rule, a.k.a. **Backpropagation**! That is the procedure to compute gradients of the loss w.r.t. parameters in a multi-layer neural network.



# Key Idea: Wiggle to Decrease Loss



- Let's say we want to decrease the loss by adjusting  $W^1_{i,j}$ .
- We could consider a very small  $\epsilon=1e-6$  and compute:

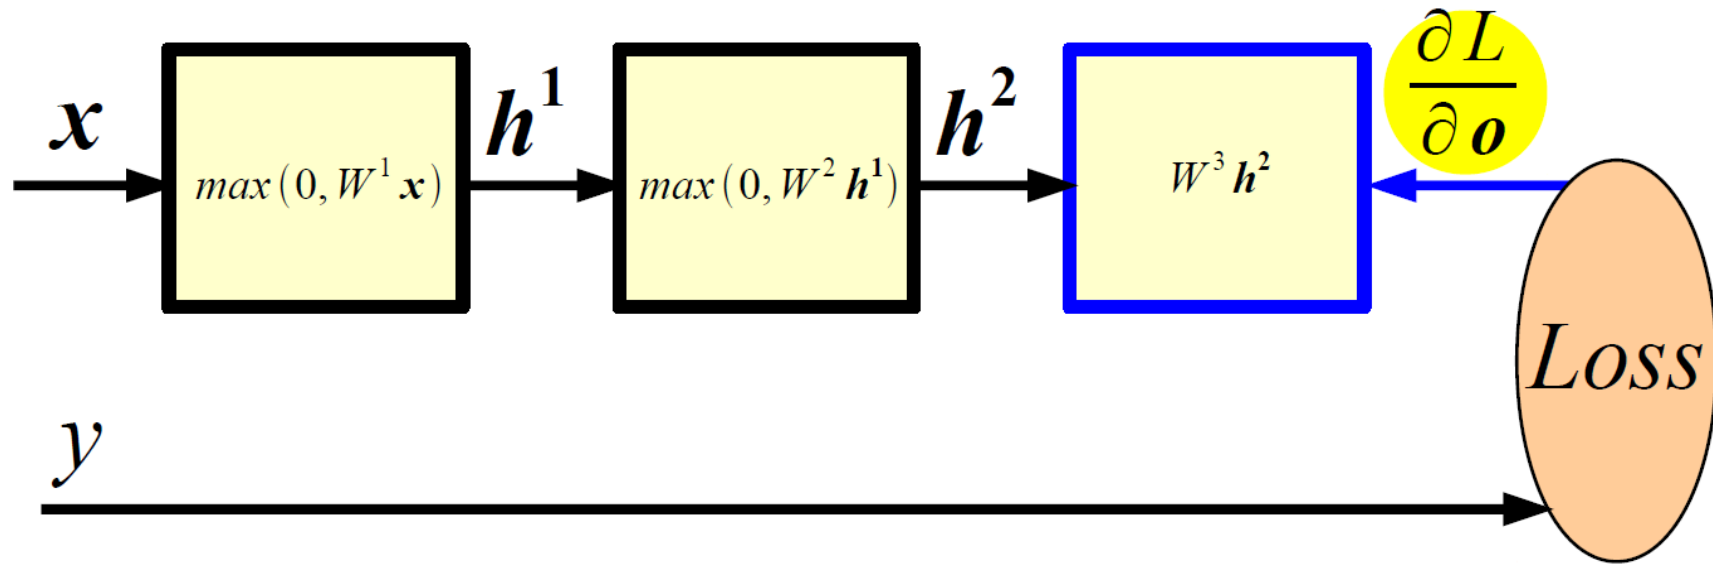
$$L(\mathbf{x}, y; \boldsymbol{\theta})$$

$$L(\mathbf{x}, y; \boldsymbol{\theta} \setminus W^1_{i,j}, W^1_{i,j} + \epsilon)$$

- Then update:

$$W^1_{i,j} \leftarrow W^1_{i,j} + \epsilon \operatorname{sgn}(L(\mathbf{x}, y; \boldsymbol{\theta}) - L(\mathbf{x}, y; \boldsymbol{\theta} \setminus W^1_{i,j}, W^1_{i,j} + \epsilon))$$

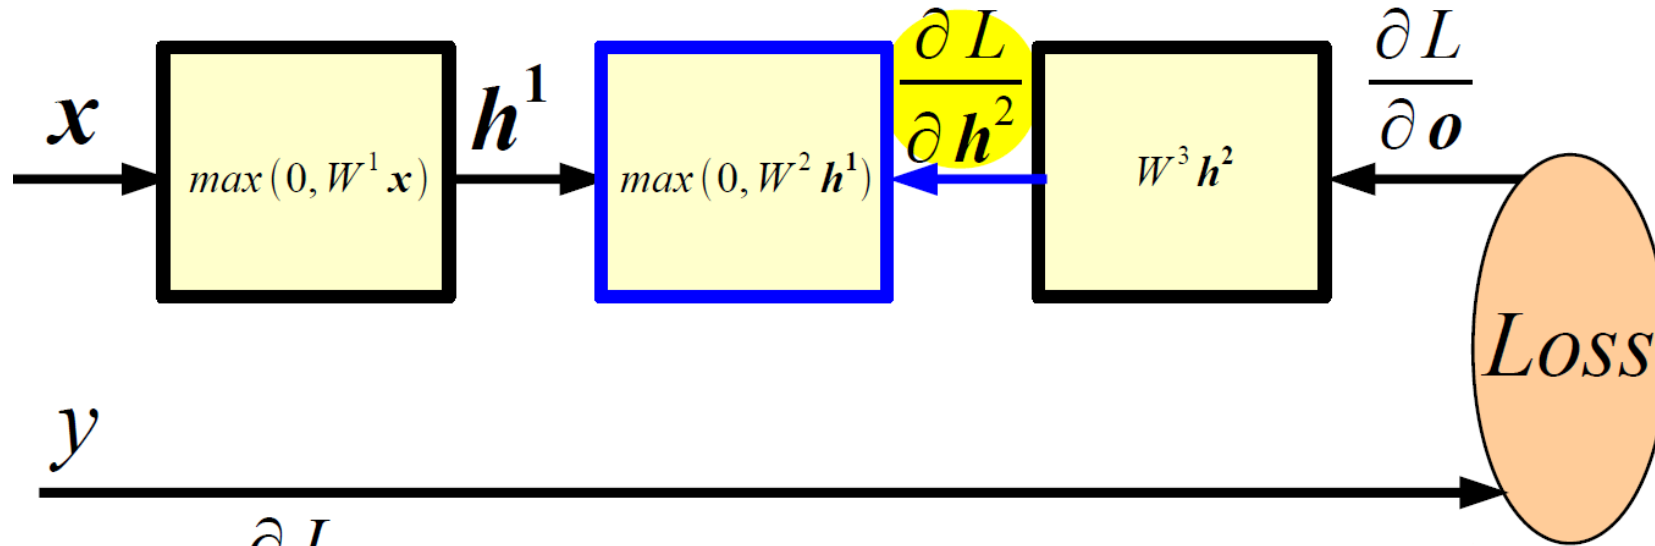
# Backward Propagation



$$\frac{\partial L}{\partial W^3} = \frac{\partial L}{\partial \mathbf{o}} \frac{\partial \mathbf{o}}{\partial W^3}$$

$$\frac{\partial L}{\partial \mathbf{h}^2} = \frac{\partial L}{\partial \mathbf{o}} \frac{\partial \mathbf{o}}{\partial \mathbf{h}^2}$$

# Backward Propagation




Given  $\frac{\partial L}{\partial h^2}$  we can compute now:

$$\frac{\partial L}{\partial W^2} = \frac{\partial L}{\partial h^2} \frac{\partial h^2}{\partial W^2}$$

$$\frac{\partial L}{\partial h^1} = \frac{\partial L}{\partial h^2} \frac{\partial h^2}{\partial h^1}$$

# Backward Propagation



Given  $\frac{\partial L}{\partial h^1}$  we can compute now:

$$\frac{\partial L}{\partial W^1} = \frac{\partial L}{\partial h^1} \frac{\partial h^1}{\partial W^1}$$

# Optimization

## Stochastic Gradient Descent

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} - \eta \frac{\partial L}{\partial \boldsymbol{\theta}}, \eta \in (0, 1)$$

Or one of its many variants